

Day 17

=====

Kubernetes

=====

Minions: This is an individual node used in kubernetes
Combination of these minions is called as Kubernetes cluster

Master is the main machine which triggers the container
orchestration

It distributes the work load to the Slaves

Slaves are the nodes that accept the work load from the
master

and handle activities load balancing, autoscaling, high
availability etc

Kubernetes uses various types of Object

1 Pod: This is a layer of abstraction on top of a
container. This is the smallest
object that kubernetes can work on. In the Pod we have a
container.

The advantage of using a Pod is that kubectl commands will
work on the Pod and the

Pod communicates these instructions to the container. In
this way we can use the

same kubectl irrespective of which technology containers
are in the Pod.

2 Service: This is used for port mapping and network load
balancing

3 Namespace: This is used for creating partitions in the cluster. Pods running in a namespace cannot communicate with other pods running in other namespace

4 Secrets: This is used for passing encrypted data to the Pods

5 ReplicationController: This is used for managing multiple replicas of PODs and also performing scaling

6 ReplicaSet: This is similar to replicationcontroller but it is more advanced where features like selector can be implemented

7 Deployment: This used for performing all activities that a Replicaset can do it can also handle rolling update

8 PersistentVolume: Used to specify the section of storage that should be used for volumes

9 PersistentVolumeClaims: Used to reserve a certain amount of storage for a pod from the persistent volume.

10 Statefulsets: These are used to handle stateful application like data bases where consistency in read write operations has to be maintained.

11 HorizontalPodAutScaler: Used for auto scaling of pods depending on the load

Kubernetes Architecture

=====

Master Componentes

=====

Container runtime: This can be docker or anyother container technology

apiServer: Users interact with the apiServer using some clinet like ui,command line tool like kubelet.It is the apiServer which is the gateway to the cluster
It works as a gatekeeper for authentication and it validates if a specific user is having permissions to execute a specific command.Example if we want to deploy a pod or a deployment first apiServers validates if the user is authorised to perform that action and if so it passes to the next process ie the "Scheduler"

Scheduler: This process accepts the instructions from apiServer after validation and starts an application on a sepcific node or set of nodes.It estimates how much amount of h/w is required for an application and then checks which slave have the necessary h/w resources and instructs the kubelet to deploy the application

kubelet: This is the actual process that takes the orders from scheduler and deploy an application on a slave.This kubelet is present on both master and slave

controller manager: This check if the desired state of the cluster is always maintained.If a pod dies it recreates that pod to maintain the desired state

etcd: Here the cluster state is maintained in key value pairs.
It maintains info about the slaves and the h/w resources available on the slaves and also the pods running on the slaves
The scheduler and the control manager read the info from this etcd and schedule the pods and maintain the desired state

=====

Worker components

=====

containerrun time: Docker or some other container technology

kubelet: This process interacts with container run time and the node
and it start a pod with a container in it

kubeproxy: This will take the request from services to pod
It has the intellegence to forward a request to a near by pod.Eg If an application pod wants to communicate with a db pod
then kubeproxy will take that request to the nearby pod

=====

Day 18

=====

Kubernetes can be installed in the following ways
Unmanaged K8s setup

=====

1 Kops

2 Kubeadm

3 Kind

Managed K8s setup

=====

1 EKS (AWS)

2 GKE (GCP)

=====

=====

KOPS stands for Kubernetes Operations and it is used for setting up Kubernetes on cloud in an automated manner

Kubernetes on AWS using Kops

1. Launch Linux EC2 instance in AWS (Kubernetes Client)
2. Create and attach IAM role to EC2 Instance.

Kops need permissions to access

- S3
- EC2
- VPC
- Route53
- Autoscaling
- etc..

3. Install Kops on EC2

```
curl -LO
```

```
https://github.com/kubernetes/kops/releases/download/$(curl -s
```

```
https://api.github.com/repos/kubernetes/kops/releases/latest | grep tag_name | cut -d '"' -f 4)/kops-linux-amd64
```

```
chmod +x kops-linux-amd64
```

```
sudo mv kops-linux-amd64 /usr/local/bin/kops
```

4. Install kubectl

```
curl -LO
https://storage.googleapis.com/kubernetes-release/release/$(
curl -s
https://storage.googleapis.com/kubernetes-release/release/st
able.txt)/bin/linux/amd64/kubect1
chmod +x ./kubect1
sudo mv ./kubect1 /usr/local/bin/kubect1
```

5. Create S3 bucket in AWS

S3 bucket is used by kubernetes to persist cluster state, lets create s3 bucket using aws cli Note: Make sure you choose bucket name that is unique accross all aws accounts

```
aws s3 mb s3://project.in.k8s --region us-west-2
```

6. Create private hosted zone in AWS Route53

Head over to aws Route53 and create hostedzone

Choose name for example (sai.in)

Choose type as privated hosted zone for VPC

Select default vpc in the region you are setting up your cluster

Hit create

7 Configure environment variables.

Open .bashrc file

```
vi ~/.bashrc
```

Add following content into .bashrc, you can choose any arbitrary name for cluster and make sure buck name matches the one you created in previous step.

```
export KOPS_CLUSTER_NAME=project.in
```

```
export KOPS_STATE_STORE=s3://project.in.k8s
```

Then running command to reflect variables added to .bashrc

```
source ~/.bashrc
```

8. Create ssh key pair

This keypair is used for ssh into kubernetes cluster

```
ssh-keygen
```

9. Create a Kubernetes cluster definition.

```
kops create cluster \  
--state=${KOPS_STATE_STORE} \  
--node-count=2 \  
--master-size=t3.medium \  
--node-size=t3.medium \  
--zones=us-west-2a \  
--name=${KOPS_CLUSTER_NAME} \  
--dns private \  
--master-count 1
```

10. Create kubernetes cluster

```
kops update cluster --yes --admin
```

Above command may take some time to create the required infrastructure resources on AWS. Execute the validate command to check its status and wait until the cluster becomes ready

11 To check if the cluster is ready

```
kops validate cluster
```

For the above above command, you might see validation failed error initially when you create cluster and it is expected behaviour, you have to wait for some more time and check again.

```
=====
=====
=KIND :Kubernetes in Docker
=====
Here the master and slave machines are docker containers
=====
===
```

1 Create a AWS ubuntu instance and install docker on it

2 Install Kubectl

```
curl -LO
```

```
https://storage.googleapis.com/kubernetes-release/release/\$\(
```

```
curl -s
https://storage.googleapis.com/kubernetes-release/release/st
able.txt)/bin/linux/amd64/kubect1
```

```
chmod +x ./kubect1
```

```
sudo mv ./kubect1 /usr/local/bin/kubect1
```

3 Install KIND

```
[ $(uname -m) = x86_64 ] && curl -Lo ./kind
https://kind.sigs.k8s.io/dl/v0.20.0/kind-linux-amd64
```

```
chmod +x ./kind
```

```
sudo mv ./kind /usr/local/bin/kind
```

4 Create a kind config file

```
vim config.yml
# three node (two workers) cluster config
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
- role: control-plane
- role: worker
- role: worker
```

5 Create the cluster with the above file

```
kind create cluster --name mycluster --config=config.yml
```

6 Cluster will be created as docker containers

```
sudo docker container ls
```

7 Go into the master container to run the kubernetes commands

```
sudo docker exec -it master_container_id bash
```


8 Run any kubernetes command
kubect1 get nodes

=====
=====
Day 19
=====
=====

Kubeadm installation-This is a manual setup for Kubernetes
and it works on both cloud and on premise
=====

=====
Install, start and enable docker service

```
yum install -y -q yum-utils device-mapper-persistent-data  
lvm2 > /dev/null 2>&1  
yum-config-manager --add-repo  
https://download.docker.com/linux/centos/docker-ce.repo >  
/dev/null 2>&1  
yum install -y -q docker-ce >/dev/null 2>&1
```

```
systemctl start docker  
systemctl enable docker
```

=====
=====
Disable SELINUX

```
setenforce 0  
sed -i --follow-symlinks  
's/^SELINUX=enforcing/SELINUX=disabled/'  
/etc/sysconfig/selinux
```

=====
=====
Disable SWAP

```
sed -i '/swap/d' /etc/fstab
swapoff -a
```

```
=====
=====
```

Update sysctl settings for Kubernetes networking

```
cat >>/etc/sysctl.d/kubernetes.conf<<EOF
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
```

```
sysctl --system
```

```
=====
=====
```

Add Kubernetes to yum repository

```
cat >>/etc/yum.repos.d/kubernetes.repo<<EOF
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubern
es-el7-x86_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
https://packages.cloud.google.com/yum/doc/rpm-package-key.gp
g
EOF
```

```
=====
=====
```

Install Kubernetes

```
yum install -y kubeadm-1.19.1 kubelet-1.19.1 kubect1-1.19.1
```

```
=====
=====
```

Enable and start Kubernetes service

```
systemctl start kubelet
systemctl enable kubelet
```

```
=====
=====
```

Repeat the above steps on Master and slaves

```
=====
=====
```

On Master=====

```
=====
```

Initilise the Kubernetes cluster

```
-----
```

```
kubeadm init --apiserver-advertise-address=ip_of_master
--pod-network-cidr=192.168.0.0/16
```

```
=====
=====
```

To be able to use kubectl command to connect and interact with the cluster,
the user needs kube config file.

```
mkdir /home/ec2-user/.kube
cp /etc/kubernetes/admin.conf /home/ec2-user/.kube/config
chown -R ec2-user:ec2-user /home/ec2-user/.kube
```

```
=====
=====
```

Deploy calico network

```
kubectl apply -f
```

```
https://docs.projectcalico.org/v3.9/manifests/calico.yaml
```

```
=====
=====
```

```
For slaves to join the cluster
kubeadm token create --print-join-command
```

```
=====
=====
```

Managed Kubernetes Installation

```
=====
  EKS (Elastic Kubernetes Service)
=====
```

1 Create an ubuntu instance on AWS and name it EKS server

2 Create an IAM with admin roles and assign to the EKS server

3 Install Kubectl

```
curl -LO
https://storage.googleapis.com/kubernetes-release/release/$(
curl -s
https://storage.googleapis.com/kubernetes-release/release/st
able.txt)/bin/linux/amd64/kubectl
```

```
chmod +x ./kubectl
```

```
sudo mv ./kubectl /usr/local/bin/kubectl
```

4 Install eksctl

Download the eksctl

curl --silent --location

```
"https://github.com/weaveworks/eksctl/releases/latest/downlo
ad/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp
```

Give execute permissions on it
sudo mv /tmp/eksctl /usr/local/bin

Check if it is installed
eksctl version

5 To create cluster on EKS
eksctl create cluster \
 --region us-east-1 \
 --node-type t3.medium \
 --nodes 3 \
 --name mynew-cluster

=====

Day 20

=====

Kubernetes Setup on GCP using GKE

=====

1 Login into GCP console

2 Click on Navigation menu

3 Click on Kubernetes Engine

4 Click on Create cluster

5 Select Switch to Standard cluster

6 Click on Create

=====

=====

To see the list of nodes in the Kubernetes cluster
kubectl get nodes

2 To get info about the nodes along with ipaddress and docker version etc

```
kubectl get nodes -o wide
```

3 To get detailed info about the nodes

```
kubectl describe nodes node_name
```

```
=====
```

Create nginx as a pod and name it webserver

```
kubectl run --image nginx webserver
```

To see the list of pods

```
kubectl get pods
```

To get info about the pods along with ipaddress

```
kubectl get pods -o wide
```

To get detailed info about the pods

```
kubectl describe pods webserver
```

```
=====
```

Create a mysql pod and also pass the necessary environment variables

```
kubectl run --image mysql:5 db --env
```

```
MYSQL_ROOT_PASSWORD=intelliqit
```

Check if the pod is running

```
kubectl get pods
```

To delete the mysql pod

```
kubectl delete pods db
```

```
=====
```

Kubernetes objects are created using definition/manifest

files

These files contain mainly four components

apiVersion:

kind:

metadata:

spec:

...

kind : apiversion

=====

Pod	v1
Service	v1
Namespace	v1
Secret	v1
ReplicationController	v1
ReplicaSet	apps/v1
Deployment	apps/v1
StatefulSet	apps/v1
DaemonSet	apps/v1
PersistentVolume	v1
PersistentVolumeClaim	v1
HorizontalPodAutoscaler	v1

=====

=====

Create a pod definition file to create an nginx pod

1 vim pod-definition1.yml

--

apiVersion: v1

kind: Pod

metadata:

name: nginx-pod

namespace: test-ns

labels:

author: intelliqit

type: proxy

```
    cat: rat
spec:
  containers:
    - name: mynginx
      image: nginx
  ...
```

2 To create pods from the above file
kubect1 apply -f pod-defintion2.yml

3 To see the list of pods
kubect1 get pods

4 To delete the pods created from theabove file
kubect1 delete -f pod-defintion1.yml

=====

Create a pod definition file to setup a postgres pod
1 vim pod-definition.yml

```
---
apiVersion: v1
kind: Pod
metadata:
  name: postgres-pod
  labels:
    type: db
    author: intelliqit
spec:
  containers:
    - name: mydb
      image: postgres
      env:
        - name: POSTGRES_PASSWORD
          value: intelliqit
        - name: POSTGRES_DB
```



```
    value: mydb
- name: POSTGRES_USER
  value: myuser
```

```
=====
Create a pod definition file to create a jenkins pod
```

```
1 vim pod-definition3.yml
```

```
---
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: jenkins-pod
```

```
  labels:
```

```
    type: ci-cd
```

```
    author: intelliqit
```

```
spec:
```

```
  containers:
```

```
    - name: myjenkins
```

```
      image: jenkins/jenkins
```

```
      ports:
```

```
        - containerPort: 8080
```

```
          hostPort: 8080
```

```
=====
Day 21
```

```
=====
Create a httpd pod using a definition file
```

```
vim pod-definition4.yml
```

```
---
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: httpd-pod
```

```
  labels:
```

```
    type: webserver
    author: intelliqit
spec:
  containers:
    - name: myhttpd
      image: httpd
      ports:
        - containerPort: 80
          hostPort: 8080
  ...
```

```
=====
=====
```

Namespace: These are logical partitions in the Kubernetes cluster

Create a definition file to create a namespace

```
vim namespace.yml
```

```
---
```

```
apiVersion: v1
kind: Namespace
metadata:
  name: test-ns
...
```

To create a namespace from the above file

```
kubectl apply -f namespace.yml
```

To see the list of all the namespaces

```
kubectl get namespace
```

Create a definition file to create wordpress and launch it on the above namespace

```
vim pod-definition5.yml
```

```
---
```

```
apiVersion: v1
kind: Pod
```

```

metadata:
  name: wordpress-pod
  namespace: test-ns
  labels:
    type: CMS
    author: intelliqit
spec:
  containers:
    - name: mywordpress
      image: wordpress
      ports:
        - containerPort: 80
          hostPort: 8080
  ...

```

To create pods from the above file
 kubectl apply -f pod-definition5.yml

To check if the pod is created on the above namespace
 kubectl get pods -n test-ns

```

=====
=====

```

ReplicationController

```
=====
```

Create a replication controller file to setup httpd with
 multiple replicas

vim replication-controller.yml

```
---
```

```

apiVersion: v1
kind: ReplicationController
metadata:
  name: httpd-rc
  labels:
    type: websrver
    author: intelliqit

```

```
spec:
  replicas: 3
  template:
    metadata:
      name: httpd-pod
      labels:
        type: webserver
    spec:
      containers:
        - name: myhttpd
          image: httpd
          ports:
            - containerPort: 80
              hostPort: 8080
  ...
```

To create replication controller from the above file
 kubectl apply -f replication-controller.yml

To see the list of replication controllers
 kubectl get rc

To see the pods
 kubectl get pods

```
=====
=====
```

ReplicaSet

```
=====
```

Create a replicaset to setup multiple replicas of tomcat
 vim replicas-set.yml

```
---
```

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: tomcat-rs
  labels:
```

```

    type: appserver
    author: intelliqit
spec:
  replicas: 3
  selector:
    matchLabels:
      type: appserver
  template:
    metadata:
      name: tomcat-pod
      labels:
        type: appserver
    spec:
      containers:
        - name: mytomcat
          image: tomee
          ports:
            - containerPort: 8080
              hostPort: 9090
...

```

To create a replicaset from the above file
 kubectl apply -f replica-set.yml

To see the list of replicaset
 kubectl get rs

To scale the replicas set we can change the no of replicas
 in the definition file and
 kubectl replace -f replicas-set.yml

Another way of scaling directly from command prompt is
 kubectl scale --replicas 1 -f replica-set.yml

```

=====
=====

```

Deployment

=====

Create a deployment definition file for nginx

vim deployment1.yml

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-deployment

labels:

type: proxy

author: intelliqit

spec:

replicas: 3

selector:

matchLabels:

type: proxy

template:

metadata:

name: nginx-pod

labels:

type: proxy

spec:

containers:

- name: mynginx

image: nginx

ports:

- containerPort: 80

hostPort: 9090

To create a deployment from the above file

kubectl apply -f deployment1.yml

To see the list of deployments

kubectl get deployment

To delete the deployments

kubectl delete -f deployment1.yml

```

=====
=====
Create a deployment definition file to setup mysql
vim deployment2.yml
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mysql-deployment
  labels:
    type: db
spec:
  replicas: 2
  selector:
    matchLabels:
      type: db
  template:
    metadata:
      name: mysql-pod
      labels:
        type: db
    spec:
      containers:
        - name: mydb
          image: mysql:5
          env:
            - name: MYSQL_ROOT_PASSWORD
              value: intelliqit
...

```

```

=====
=====
Day 22
=====
=====
DaemonSet: This is to run pods on every salve and only one

```

slave per node.

Create a daemonset file to create nginx

vim daemonset.yml

apiVersion: apps/v1

kind: DaemonSet

metadata:

name: nginx-daemon

labels:

type: proxy

spec:

selector:

matchLabels:

type: proxy

template:

metadata:

name: nginx-pod

labels:

type: proxy

spec:

containers:

- name: mynginx

image: nginx

ports:

- containerPort: 80

hostPort: 8080

...

To create a daemonset from this file

kubectl apply -f daemonset.yml

To see the list of all the objects running in the cluster

kubect get all

=====
=====

Service Objects

=====

1 NodePort: This is used to perform network load balancing

2 LoadBalancer: This will create an ip for the entire cluster and it works only on managed kubernetes service

3 Clusterip: This is used to fro pods to communicate with other pods in the clsuter but not with outside world

=====

=====

Create a service definition file fro node port object and apply it on pod-definition1.yml

```
vim service1.yml
---
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
  labels:
    author: intelliqit
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
    type: proxy
```

Create the pod for nginx using the definition file

kubect1 apply -f pod-definition1.yml

Create service from the above file

kubect1 apply -f service1.yml

To see the list of all object
kubectl get all

Now we can access nginx from any machines ip address

=====

Create a service object of the type loadbalancer and apply
it on pod-definition3.yml

vim service2.yml

apiVersion: v1

kind: Service

metadata:

name: jenkins-service

labels:

author: intelliqit

spec:

type: LoadBalancer

ports:

- targetPort: 8080

port: 8080

nodePort: 30009

selector:

type: ci-cd

author: intelliqit

...

Create a jenkins pod

kubectl apply -f pod-definition3.yml

Create a loadbalancer service from the above file

kubectl apply -f service2.yml

This will generate a unique public ip for the entire cluster

kubectl get svc

```
=====
=====
```

Create a service object of the type clusterip and apply it on pod-definition2.yml

```
vim service3.yml
```

```
---
```

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: postgres-service
```

```
  labels:
```

```
    author: intelliqit
```

```
spec:
```

```
  ports:
```

```
    - targetPort: 5432
```

```
      port: 5432
```

```
  selector:
```

```
    type: db
```

```
    author: intelliqit
```

Create postgres pod

```
kubectl apply -f pod-definition2.yml
```

Create service of clusterip type

```
kubectl apply -f service3.yml
```

To check if the service is created

```
kubectl get svc
```

```
=====
=====
```

Day 23

```
=====
=====
```

Kompose

This is used to convert a docker compose file to Kubernetes

definition files

Install Kompose into the Kubernetes cluster

<https://www.digitalocean.com/community/tutorials/how-to-migrate-a-docker-compose-workflow-to-kubernetes>

Create a docker compose file

```
vim docker-compose.yml
```

```
---
version: '3'
services:
  db:
    image: mysql:5
    environment:
      MYSQL_ROOT_PASSWORD: intelliqit
    deploy:
      replicas: 2

  wordpress:
    image: wordpress
    ports:
      - 8888:80
    deploy:
      replicas: 3
...
```

To create Kubernetes definition files from the above file

kompose convert

```
=====
=====
```

Kubernetes Project

```
=====
```

This is a voting app created using python, this app is exposed to the customers and they can cast their vote. This info will be registered in an in-memory db (temporary db) that we setup using redis.

From here we have a .net application that filters the data and stores it permanently in a postgres db and the results can be viewed on an app created using nodejs

Create 5 deployment definition files for all the above object and 4 service definition file

```
vim voting-app-deployment.yml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: voting-app-deployment
```

```
  labels:
```

```
    name: voting-app
```

```
    author: intelliqit
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      name: voting-app
```

```
  template:
```

```
    metadata:
```

```
      name: voting-app-pod
```

```
      labels:
```

```
        name: voting-app
```

```
    spec:
```

```
      containers:
```

```
        - name: voting-app
```

```
          image: dockersamples/examplevotingapp_vote
```

```
...
```

```
vim result-app-deployment.yml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
name: result-app-deployment
labels:
  name: result-app
  author: intelliqit
spec:
  replicas: 2
  selector:
    matchLabels:
      name: result-app
  template:
    metadata:
      name: result-app-pod
      labels:
        name: result-app
    spec:
      containers:
        - name: result-app
          image: dockersamples/examplevotingapp_result
...

```

vim redis-app-deployment.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis-app-deployment
  labels:
    name: redis-app
    author: intelliqit
spec:
  selector:
    matchLabels:
      name: redis-app
  template:
    metadata:
      name: redis-app-pod
      labels:

```

```
    name: redis-app
spec:
  containers:
    - name: redis-app
      image: redis
...

vim postgres-app-deployment.yml
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: postgres-app-deployment
  labels:
    name: postgres-app
    author: intelliqt
spec:
  selector:
    matchLabels:
      name: postgres-app
  template:
    metadata:
      name: postgres-app-pod
      labels:
        name: postgres-app
    spec:
      containers:
        - name: postgres-app
          image: postgres
          env:
            - name: POSTGRES_PASSWORD
              value: intelliqt
            - name: POSTGRES_USER
              value: myuser
            - name: POSTGRES_DB
              value: mydb
...
```

```
vim worker-app-deployment.yml
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: worker-app-deployment
  labels:
    name: worker-app
    author: intelliqit
spec:
  selector:
    matchLabels:
      name: worker-app
  template:
    metadata:
      name: worker-app-pod
      labels:
        name: worker-app
    spec:
      containers:
        - name: worker-app
          image: dockersamples/examplevotingapp_worker
...

```

```
vim voting-app-service.yml
---
apiVersion: v1
kind: Service
metadata:
  name: voting-app-service
  labels:
    author: intelliqit
spec:
  type: LoadBalancer
  ports:
    - targetPort: 80

```



```
    port: 80
    nodePort: 30008
  selector:
    name: voting-app
```

```
vim result-app-service.yml
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: result-app-service
  labels:
    author: intelliqit
spec:
  type: LoadBalancer
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30009
  selector:
    name: result-app
```

```
vim redis-app-service.yml
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: redis-app-service
  labels:
    author: intelliqit
spec:
  ports:
    - targetPort: 6379
      port: 6379
  selector:
    name: redis-app
```

```
...
```

```
vim postgres-app-service.yml
```

```
---
```

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: postgres-app-service
```

```
  labels:
```

```
    author: intelliqit
```

```
spec:
```

```
  ports:
```

```
    - targetPort: 5432
```

```
      port: 5432
```

```
  selector:
```

```
    name: postgres-app
```

```
...
```

```
kubectl apply -f voting-app-deployment.yml
```

```
kubectl apply -f voting-app-service.yml
```

```
kubectl apply -f result-app-deployment.yml
```

```
kubectl apply -f result-app-service.yml
```

```
kubectl apply -f redis-app-deployment.yml
```

```
kubectl apply -f redis-app-service.yml
```

```
kubectl apply -f postgres-app-deployment.yml
```

```
kubectl apply -f postgres-app-service.yml
```

```
kubectl apply -f worker-app-deployment.yml
```

```
=====
```

```
=====
```

```
Day 24 Session-1
```

```
=====
```

```
=====
```

```
Requests and Limits
```

```
=====
```

```
This is used to specify the min and max amount of hardware  
to be allocated for the pods
```

Request is used to specify the minimum amount of hardware
and limits is used to specify the
maximum amount of hardware

```
vim RequestsandLimits1.yml
```

```
---
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    author: intelliqit
    type: proxy
spec:
  containers:
  - name: mynginx
    image: nginx
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
...
```

To create pod from the above file
`kubectl apply -f RequestsandLimits1.yml`

To check the hardware allocated for the pod
`kubectl describe pods nginx-pod`
Here we will see the min and max amount of hardware

```
=====
=====
vim RequestsandLimits2.yml
```

```

---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    type: proxy
spec:
  replicas: 2
  selector:
    matchLabels:
      type: proxy
  template:
    metadata:
      name: nginx-pod
      labels:
        type: proxy
    spec:
      containers:
        - name: mynginx
          image: nginx
          resources:
            requests:
              cpu: "250m"
              memory: "64Mi"
            limits:
              cpu: "500m"
              memory: "128Mi"

```

To create deployment from the above file
 kubectl apply -f RequestsandLimits2.yml

To check the properties of the above deployment
 kubectl describe deployment nginx-deployment

```

=====
=====

```

Secrets

=====

These are Kubernetes objects that are used to pass confidential info to the pods like usernames and passwords etc

```
vim Secret1.yml
```

```
---
```

```
apiVersion: v1
kind: Secret
metadata:
  name: mysql-secret
  labels:
    author: intelliqit
type: Opaque
stringData:
  a: aW50ZWxsaXFpdAo=
...
```

To create secret from the above
`kubectl apply -f Secret1.yml`

```
vim pod-definition6.yml
```

```
---
```

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql-pod
  labels:
    type: db
spec:
  containers:
    - name: mydb
      image: mysql:5
      env:
        - name: MYSQL_ROOT_PASSWORD
```

```
    valueFrom:
      secretKeyRef:
        name: mysql-secret
        key: a
...

```

Create pod using the above file
kubectl apply -f pod-definition6.yml

```
=====
=====
vim Secret2.yml

```

```
---
apiVersion: v1
kind: Secret
metadata:
  name: postgres-secret
  labels:
    author: intelliqit
type: Opaque
stringData:
  password: intelliqit
  username: myuser
  dbname: mydb
...

```

Create a secret from the above file
kubectl apply -f Secret2.yml

Create a deployment file to use the above secret
vim deployment3.yml

```
---
apiVersion: apps/v1
kind: Deployment
metadata:

```

```

name: postgres-deployment
labels:
  type: db
spec:
  replicas: 2
  selector:
    matchLabels:
      type: db
  template:
    metadata:
      name: postgres-pod
      labels:
        type: db
    spec:
      containers:
        - name: mydb
          image: postgres
          env:
            - name: POSTGRES_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: postgres-secret
                  key: password
            - name: POSTGRES_USER
              valueFrom:
                secretKeyRef:
                  name: postgres-secret
                  key: username
            - name: POSTGRES_DB
              valueFrom:
                secretKeyRef:
                  name: postgres-secret
                  key: dbname
...

```

Create a deployment from the above file
 kubectl apply -f deployment3.yml

```
=====
NodeAffinity
=====
```

This feature is used to ensure that pods are attracted to a specific slave/node

To perform this we should label the node

```
kubectl label nodes node_name slave1=intelliqit1
```

```
vim NodeAffinity1.yml
```

```
---
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    type: proxy
spec:
  containers:
    - name: mynginx
      image: nginx
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: slave1
                operator: In
                values:
                  - intelliqit1
...

```

To create the pod

```
kubectl apply -f NodeAffinity1.yml
```

Check where the pod is running


```
kubect1 get pods -o wide
```

The pod should run only on the above labeled node

```
=====
=====
```

Day 24 Session 2

```
=====
=====
```

```
vim NodeAffinity2.yml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nginx-deployment
```

```
  labels:
```

```
    type: proxy
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      type: proxy
```

```
  template:
```

```
    metadata:
```

```
      name: nginx-pod
```

```
      labels:
```

```
        type: proxy
```

```
    spec:
```

```
      containers:
```

```
        - name: mynginx
```

```
          image: nginx
```

```
      affinity:
```

```
        nodeAffinity:
```

```
          requiredDuringSchedulingIgnoredDuringExecution:
```

```
            nodeSelectorTerms:
```

```
              - matchExpressions:
```

```
                - key: slave1
```

```
operator: In
values:
  - intelliqit1
```

...

Create deployment from the above file
kubectl apply -f NodeAffinity2.yml

Check where the pods are running
kubectl get pods -o wide

```
=====
=====
```

Taints and Tolerations

```
=====
```

Taint is used to repel pods from a specific slave

Taint is applied at the level of the slave and pods will not run on the tainted machine

```
kubectl taint nodes node_name slave1=intelliqit1:NoSchedule
```

Tolerations are used if we want to run a pod on a tainted machine

```
vim Tolerations1.yml
```

```
---
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: httpd-pod
```

```
  labels:
```

```
    type: webserver
```

```
spec:
```

```
  containers:
```

```
    - name: myhttpd
```

```
      image: httpd
```

```
  tolerations:
```

```
    - key: slave2
```

```
      operator: Equal
```

```
      value: intelliqit2
```

```
    effect: NoSchedule
```

```
...
```

Create pod from the above file and it will run on the tainted machine

```
kubectl apply -f Tolerations1.yml
```

Check where the pod is running

```
kubectl get pods -o wide
```

```
=====
```

```
vim Tolerations2.yml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: httpd-deployment
```

```
  labels:
```

```
    type: webserver
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      type: webserver
```

```
  template:
```

```
    metadata:
```

```
      name: httpd-pod
```

```
      labels:
```

```
        type: webserver
```

```
    spec:
```

```
      containers:
```

```
        - name: myhttpd
```

```
          image: httpd
```

```
      tolerations:
```

```
        - key: slave2
```

```
operator: Equal
value: intelliqit2
effect: NoSchedule
```

...

To create a deployment from the above file
kubectl apply -f Tolerations2.yml

Check where the pods are running
kubectl get pods -o wide

```
=====
=====
```

Volumes

```
=====
```

This is used for preserving the data even if the pods are crashed

vim Volumes1.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: redis-pod
  labels:
    type: db
spec:
  containers:
    - name: myredis
      image: redis
      volumeMounts:
        - name: myvolume
          mountPath: /data/redis
  volumes:
    - name: myvolume
      emptyDir: {}
```

...

Create a pod from the above file
kubectl apply -f Volumes1.yml

Go inside the pods
kubectl exec -it redis-pod -- bash

cd redis
Create a file here

Delete the pod
kill 1

Check if the pod has recreated
kubectl get pods

Go inside th new pod and we will see the data
kubectl exec -it redis-pod -- bash
cd redis
ls

=====

Day 25

=====

Rolling updates

=====

This feature is used to update from one version to another
without downtime

Create deployment with nginx:1.24 and upgrade to nginx:1.25

vim deployment1.yml

apiVersion: apps/v1
kind: Deployment
metadata:

```

name: nginx-deployment
labels:
  type: proxy
  author: intelliqit
spec:
  replicas: 2

  selector:
    matchLabels:
      type: proxy
  template:
    metadata:
      name: nginx-pod
      labels:
        type: proxy
    spec:
      containers:
        - name: nginx
          image: nginx:1.24
          ports:
            - containerPort: 80
              hostPort: 9090

```

Create deployment from the above file
 kubectl apply -f deployment1.yml

Check which version it is working on
 kubectl describe deployment nginx-deployment | less

Upgrade to nginx:1.25 version
 kubectl set image deployment/nginx-deployment
 nginx=nginx:1.16.1

Check which version it is working on
 kubectl describe deployment nginx-deployment | less

```

=====
=====

```

Recreate update strategy

=====

In the scenario it initially remove the older versions and replaces it with new versions

```
vim deployment1.yml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nginx-deployment
```

```
  labels:
```

```
    type: proxy
```

```
    author: intelliqit
```

```
spec:
```

```
  replicas: 2
```

```
  strategy:
```

```
    type: Recreate
```

```
  selector:
```

```
    matchLabels:
```

```
      type: proxy
```

```
  template:
```

```
    metadata:
```

```
      name: nginx-pod
```

```
      labels:
```

```
        type: proxy
```

```
    spec:
```

```
      containers:
```

```
        - name: nginx
```

```
          image: nginx:1.24
```

```
          ports:
```

```
            - containerPort: 80
```

```
              hostPort: 9090
```

```
...
```

Create a deployment from the above file

```
kubectl apply -f deployment1.yml
```

```
Upgrade to higher version
kubectl set image deployment/nginx-deployment
nginx=nginx:1.16.1
```

```
=====
=====
```

Blue green deployments

```
=====
```

Here the current version is called blue deployment and the latest version is called green deployments
Both of them are maintained and once the latest version is exposed to the client based on the feedback we will remove the older version

```
vim blue-deployment.yml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: blue-nginx-deployment
```

```
  labels:
```

```
    type: proxy
```

```
    author: intelliqit
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      type: proxy
```

```
  template:
```

```
    metadata:
```

```
      name: nginx-pod
```

```
      labels:
```

```
        type: proxy
```

```
    spec:
```

```
      containers:
```

```
        - name: nginx
```



```
image: nginx:1.24
ports:
  - containerPort: 80
    hostPort: 9090
```

Setup the blue deployments
kubectl apply -f blue-deployment.yml

Create a deployment file for the latest version
vim green-deployment.yml

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: green-nginx-deployment
  labels:
    type: proxy
    author: intelliqit
spec:
  replicas: 2
  selector:
    matchLabels:
      type: proxy
  template:
    metadata:
      name: nginx-pod
      labels:
        type: proxy
    spec:
      containers:
        - name: nginx
          image: nginx:1.25
          ports:
            - containerPort: 80
              hostPort: 9090
```

Create green deployment from the above file

```
kubectl apply -f green-deployment.yml
```

Delete the blue deployment.yml

```
kubectl delete -f blue-deployment.yml
```

```
=====
```

Canary deployments

```
=====
```

Here the older version is scaled down incrementally and the latest version is scaled up incrementally

```
=====
```

AutoScaling

```
=====
```

This is used to automatically scale the no of pods based on h/w utilization

```
vim Autoscalling.yml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: php-apache
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      run: php-apache
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        run: php-apache
```

```
    spec:
```

```
      containers:
```

```
        - name: php-apache
```

```
    image: intelliqit/mynew
    ports:
      - containerPort: 80
    resources:
      requests:
        cpu: 200m
      limits:
        cpu: 500m
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: php-apache
  labels:
    run: php-apache
spec:
  ports:
    - port: 80
  selector:
    run: php-apache
...
```

Create a deployment and service from the above file

```
kubectl apply -f Autoscalling.yml
```

Create a horizontal pod autoscaler

```
kubectl scale deployment php-apache --cpu-percent=50 --min=1
--max=10
```

Now apply load on the above deployment

```
kubectl run -i --tty lg --image=busybox -- /bin/sh -c "while
sleep 0.01; do wget -q -O- http://php-apache; done"
```

To see the pods autoscalling

```
kubectl get hpa php-apache --watch
```

=====

=====

HELM

=====

This is a package management software for kubernetes using which we can install or uninstall s/w's in the Kubernetes cluster

Helm uses charts for creating various components
helm create chartname

Create a helm chart for nginx and expose using service "LoadBalancer"
helm create mynginx

Go into the helm chart and edit the values.yml
cd mynginx
vim values.yml
Change the service from clusterip to LoadBalancer

Come out of the helm chart and create a release from the above chart
helm install nginx mynginx

Check the components that are created from the above chart
kubectl get all | less

=====

=====

Day 26

=====

=====

Download helm chart from artifacthub.io to setup wordpress in combination with mariadb

1 Go to artifact hub.io and search for wordpress

2 Click on first wordpress suggestion coming from "bitnami"

3 Copy the command to setup wordpress

```
helm install my-release
```

```
oci://registry-1.docker.io/bitnamicharts/wordpress
```

4 Check the components that are created

```
kubectl get all | less
```

```
=====
=====
```

Prometheus grafana

```
=====
```

Prometheus is a monitoring tool using which we can capture metrics from Kubernetes cluster

Grafana is a dashbaord for graphical visualizations of the data coming from prometheus

Install prometheus and grafana using helm

```
=====
```

```
helm repo add prometheus-community
```

```
https://prometheus-community.github.io/helm-charts
```

```
helm repo add stable https://charts.helm.sh/stable
```

```
helm repo update
```

```
helm install prometheus
```

```
prometheus-community/kube-prometheus-stack
```

Grafana by default runs on clusterip to make to accessable externally change to nodeport or loadbalancer

```
kubectl patch svc prometheus-grafana -p '{"spec": {"type": "LoadBalancer"}}'
```

Access grafana using the loadbalancer public ip in a browser

Username is admin
password: prom-operator

Grafana uses various dashboards that can be downloaded from
<https://grafana.com/grafana/dashboards/>

=====

StatefulSet

=====

A StatefulSet is a set of pods with a unique, persistent hostname and ID. StatefulSets are designed to run stateful applications in Kubernetes with dedicated persistent storage. When pods run as part of a StatefulSet, Kubernetes keeps state data in the persistent storage volumes of the StatefulSet, even if the pods shut down.

StatefulSets are commonly used to run replicated databases with a unique persistent ID for each pod.

=====

Day 27

=====

